

Chained.chat – Latest Multi-LLM Orchestration Resources (Nov 2024–Jan 2025)

Multi-Agent Orchestration Frameworks

- **CrewAI** – *GitHub*: [crewAIInc/crewAI](https://github.com/crewAIInc/crewAI) – 42k★, updated Jan 2025. A fast, **open-source Python framework for orchestrating autonomous AI agents** (“Crews”) and event-driven workflows (“Flows”), built from scratch (no LangChain dependencies) ¹. *Relevance*: CrewAI is a **widely adopted LangGraph alternative** (40k+ stars ²), offering high performance and granular control for complex multi-agent tasks in production (even used by Fortune 500 companies) ².
- **LangGraph** – *GitHub*: [langchain-ai/langgraph](https://github.com/langchain-ai/langgraph) – 23k★, updated Dec 2024. A **low-level orchestration framework to build long-running, stateful LLM agents as graphs** of nodes and edges ³. It provides durable, resumable execution, human-in-the-loop inspection, and persistent memory for agent workflows ⁴. *Relevance*: LangGraph (from the LangChain ecosystem) is a **direct competitor to CrewAI**, enabling robust multi-step agent plans and integrations (used by companies like Replit and Klarna) for resilient AI agent applications ³.
- **Microsoft AutoGen** – *GitHub*: [microsoft/autogen](https://github.com/microsoft/autogen) – ≈50k★, updated through 2024. An **event-driven multi-agent programming framework** by Microsoft for creating agents that collaborate or interface with humans. It pioneered many orchestration patterns now common in other frameworks ⁵. *Relevance*: AutoGen’s huge community (50k+ stars as of late 2025 ⁵) and enterprise use (now evolving into the unified Microsoft **Agent Framework**) make it a **reference architecture** for agent orchestration – informing best practices for building your platform’s agent system.

(Note: Temporal workflows are also being leveraged for LLM agents – e.g. OpenAI’s Codex agent runs on Temporal for durable, stateful orchestration ⁶ ⁷. This pattern shows how deterministic workflows (Temporal) can ensure reliability and recovery in dynamic AI agent processes, complementing the above frameworks.)

LLM Gateway/Proxy Solutions

- **Portkey AI Gateway** – *GitHub*: [Portkey-AI/gateway](https://github.com/Portkey-AI/gateway) – 10.2k★, updated Dec 2024. A **blazing-fast unified LLM gateway** (written in Go) with built-in **routing, load balancing, guardrails, and support for 200+ models** behind an OpenAI-compatible API ⁸. It offers org-level features (rate limiting, auth, cost controls) out-of-the-box. *Relevance*: Portkey is a popular **LiteLLM alternative** focused on performance and safety – ideal for chained.chat to manage multiple providers, enforce quotas, and add fallbacks with minimal overhead.
- **Helicone** – *GitHub*: [Helicone/helicone](https://github.com/Helicone/helicone) – 4.9k★, updated Dec 2024. An **open-source LLM observability proxy** and gateway that sits between your app and providers ⁹. It logs *every* request/response with rich metadata (latency, tokens, cost) into a sleek dashboard ⁹ ¹⁰.

Relevance: Helicone can be used alongside your API calls or other gateways to **track usage, spending, and performance** across OpenAI, Anthropic, etc. – crucial for cost monitoring and debugging in a multi-LLM platform.

- **TensorZero** – *GitHub:* [tensorzero/tensorzero](https://github.com/tensorzero/tensorzero) – **10.8k★, updated Jan 2025**. A **high-performance Rust-based LLM gateway** that unifies model routing with first-class **observability, evaluation, and optimization features** ¹¹ ¹². It's effectively an "all-in-one" LiteLLM alternative – <1 ms added latency, built-in token caching, A/B testing, feedback logging, etc. ¹² ¹³. *Relevance:* TensorZero's design (Rust speed + integrated evals/experiments) is **aligned with enterprise needs** – it demonstrates how to achieve low overhead and advanced monitoring, which is valuable for informing chained.chat's own gateway layer (or integrating with it).

(Other notable mentions: OpenRouter (a unified API proxy for many models, managed by Together AI) and LiteLLM itself are used in this space ¹⁴, but the projects above represent more recent open-source solutions with active development after Nov 2024.)

Semantic Caching Libraries

- **GPTCache** – *GitHub:* [zilliztech/GPTCache](https://github.com/zilliztech/GPTCache) – **7.9k★, updated through 2024**. A library for **semantic caching of LLM query responses** using embeddings and vector similarity ¹⁵. By storing question-answer pairs as vectors, GPTCache can reuse results for semantically similar queries, **cutting API costs by up to 10× and latency by 100×** ¹⁶. *Relevance:* Integrating a semantic cache like GPTCache into chained.chat could **dramatically reduce redundant LLM calls** – it's already integrated with LangChain and LlamaIndex ¹⁷, showing it can plug into LLM workflows to save money and time.
- **Codefuse ModelCache** – *GitHub:* [codefuse-ai/ModelCache](https://github.com/codefuse-ai/ModelCache) – **~960★, updated Oct 2024**. An **LLM semantic caching system** that stores embeddings of prompts and model outputs to quickly return results for similar queries ¹⁸. It supports both local (SQLite/FAISS) and distributed backends (MySQL+Milvus), and even added **Redis vector search** support for multi-tenant caching (Oct 2024 update) ¹⁹. *Relevance:* ModelCache is purpose-built to **speed up LLM apps by caching Q&A pairs**. Its recent addition of Redis integration and multi-language embedding support shows it's evolving with industry best practices, which could inspire the caching layer for your platform.
- **Upstash Semantic Cache (Node)** – *GitHub:* [upstash/semantic-cache](https://github.com/upstash/semantic-cache) – **290★, updated mid-2024**. A lightweight **semantic text cache for Node.js** that uses Upstash's Redis vectors under the hood. It lets you cache and retrieve values by meaning (e.g. "best places to visit in Spain" hits cache for "great places to check out in Spain") ²⁰ ²¹. *Relevance:* Demonstrates a simple fuzzy KV-store approach to semantic caching. This is relevant if chained.chat wants a **Redis-based semantic cache** – Upstash's example shows how to implement one with a vector DB, which could be replicated or integrated for the platform's caching mechanism.

(Note: Major infrastructure players are also pushing semantic caching – e.g. in April 2025 Redis Inc. launched Vector Sets and a managed LangCache service for semantic LLM caching ²² ²³. This trend underlines the importance of embedding-driven caches to optimize LLM apps.)

Real-Time Streaming UI Components

- **FlowToken (React)** – *GitHub*: [Ephibbs/flowtoken](https://github.com/eph1bbs/flowtoken) – 490★, first released Dec 2024. A React component library for smooth streaming text animations from LLMs ²⁴. It provides ready-made effects (typewriter, fade-in, slide, word-by-word reveal, etc.) to make the AI's response text appear fluidly as it streams ²⁵. *Relevance*: FlowToken can greatly **improve the UX of chained.chat's front-end**. Instead of bluntly appending tokens, you can use this to animate responses with fades or flips, creating a polished, engaging chat interface out-of-the-box.
- **"Chasing 240 FPS in LLM Chat UIs"** – *Blog post*: (Dev.to, Dec 15 2025) ²⁶ ²⁷. This in-depth article benchmarks techniques to keep a React chat UI ultra-smooth during **streaming LLM responses**. Key takeaways include building state outside React (using stores like Zustand), using windowing/virtualization instead of rendering huge conversation lists, leveraging CSS optimizations (`content-visibility: auto`, `will-change`) and even intentionally slowing the token stream slightly for performance ²⁶ ²⁷. *Relevance*: Although more of a guide than code, it provides **battle-tested patterns for streaming UI**. Applying these lessons will help chained.chat deliver **real-time results without lag** – crucial as conversations grow long or multi-agent outputs arrive concurrently.
- **Vercel AI SDK (React)** – *Docs*: [useChat – streaming example](https://sdk.vercel.ai/docs/streaming) – (N/A). Vercel's open-source AI SDK offers a `useChat` React hook that handles server-sent event streaming from LLMs ²⁸. With minimal code, it maintains a message list state and streams tokens to the UI. *Relevance*: This SDK (which underpins Vercel's own Chat demo) illustrates **best practices for streaming responses in a Next.js/React app**. You might not use it directly, but its patterns (SSE via fetch, incremental state updates, error handling) can guide the implementation of chained.chat's live response viewer.

Each of these resources was **updated after Nov 2024** and reflects the state-of-the-art. By leveraging these frameworks, gateways, caching systems, and UI components, **chained.chat** can quickly assemble a robust, efficient multi-LLM orchestration platform with proven ideas from the community. All links above point directly to the projects or references for further exploration of features and documentation.

Sources:

- Multi-Agent frameworks: CrewAI ² ¹, LangGraph ³ ⁴, Microsoft AutoGen ⁵, Temporal usage ⁶ ⁷.
- LLM Gateways: Portkey ⁸, Helicone ⁹, TensorZero ¹¹ ¹².
- Semantic Cache: GPTCache ¹⁶, ModelCache ¹⁹ ¹⁸, Upstash cache ²⁰ ²¹, Redis LangCache news ²² ²³.
- Streaming UI: FlowToken ²⁵, 240 FPS UI blog ²⁶ ²⁷, Vercel useChat hook ²⁸.

- 1 GitHub - crewAIInc/crewAI: Framework for orchestrating role-playing, autonomous AI agents. By fostering collaborative intelligence, CrewAI empowers agents to work together seamlessly, tackling complex tasks.
<https://github.com/crewAIInc/crewAI>
- 2 Top Tools for Event-Driven LLM Workflow Design
<https://latitude-blog.ghost.io/blog/top-tools-event-driven-llm-workflow-design/>
- 3 4 GitHub - langchain-ai/langgraph: Build resilient language agents as graphs.
<https://github.com/langchain-ai/langgraph>
- 5 AutoGen Update · microsoft/autogen · Discussion #7066 · GitHub
<https://github.com/microsoft/autogen/discussions/7066>
- 6 7 Of course you can build dynamic AI agents with Temporal | Temporal
<https://temporal.io/blog/of-course-you-can-build-dynamic-ai-agents-with-temporal>
- 8 GitHub - Portkey-AI/gateway: A blazing fast AI Gateway with integrated guardrails. Route to 200+ LLMs, 50+ AI Guardrails with 1 fast & friendly API.
<https://github.com/Portkey-AI/gateway>
- 9 10 Top 5 LiteLLM alternatives of 2025
<https://www.truefoundry.com/blog/litellm-alternatives>
- 11 GitHub - tensorzero/tensorzero: TensorZero is an open-source stack for industrial-grade LLM applications. It unifies an LLM gateway, observability, optimization, evaluation, and experimentation.
<https://github.com/tensorzero/tensorzero>
- 12 13 Comparison: TensorZero vs. LiteLLM - TensorZero Docs
<https://www.tensorzero.com/docs/comparison/litellm>
- 14 Top 5 LiteLLM Alternatives in 2025 - DEV Community
<https://dev.to/debmckinney/top-5-litellm-alternatives-in-2025-1pki>
- 15 16 17 zilliztech/GPTCache: Semantic cache for LLMs. Fully ... - GitHub
<https://github.com/zilliztech/GPTCache>
- 18 19 GitHub - codefuse-ai/ModelCache: A LLM semantic caching system aiming to enhance user experience by reducing response time via cached query-result pairs.
<https://github.com/codefuse-ai/ModelCache>
- 20 21 GitHub - upstash/semantic-cache: A fuzzy key value store based on semantic similarity rather lexical equality.
<https://github.com/upstash/semantic-cache>
- 22 23 Redis Vector Sets and LangCache speed GenAI models
<https://blocksandfiles.com/2025/04/11/redis-vector-sets-and-langcache-speed-genai-models/>
- 24 25 GitHub - Ephibbs/flowtoken: A UI library to animate and style streaming LLM output
<https://github.com/Ephibbs/flowtoken>
- 26 27 Chasing 240 FPS in LLM Chat UIs - DEV Community
https://dev.to/gokhan_koc_88338a026508b3/chasing-240-fps-on-llm-chats-4gde
- 28 Basic Chatbot | Vercel Academy
<https://vercel.com/academy/ai-sdk/basic-chatbot>